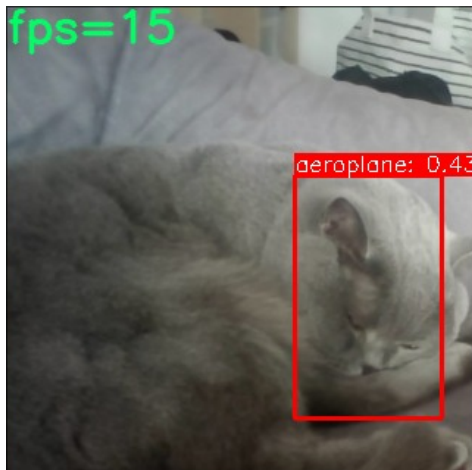


Embedded Machine Learning Lab

Max Schik, Darius Schefer | March 20, 2024

Contents

Cat



Data

More data

- Human Dataset ^a (17,300 images)
- Tiktok Dancing ^b (2615 images)

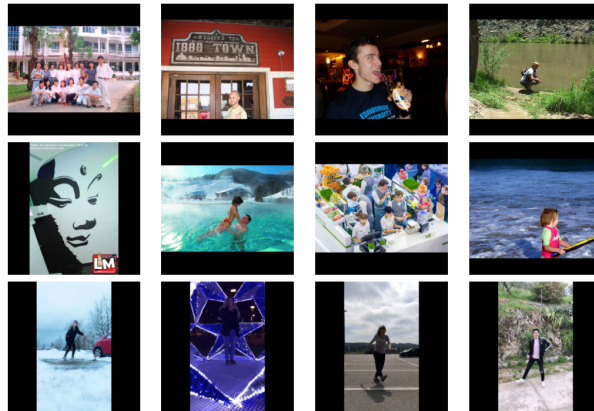
Data augmentation

- Albumentations^c library
- Rotation, flipping, contrast

^a<https://www.kaggle.com/datasets/fareselmenshawii/human-dataset>


^b<https://www.kaggle.com/datasets/tapakah68/segmentation-full-body-tiktok-dancing-dataset>




^c<https://albumentations.ai/>



Adaptation and optimization

- Person-only-detection, fine-tuning
- Iterative pruning
- Batch norm optimization
- ONNX Runtime
- TensorRT





Integration and pipeline

■ Configuration

- File containing steps and parameters
- Reproducibility
- Version control

■ Camera loop

- Jupyter Notebook
- Socket slow

```
1 --- !experiment
2 start_weights_path: "./weights/voc_pretrained.pt"
3 augment: True
4
5 steps:
6 - !strip_classes
7   finetune: true
8   finetune_epochs: 15
9 - !pruning
10   target_acc: 0.3
11   prune_ratio: 0.05
12   batch_size: 64
13   num_train_epochs: 10
14   num_eval_batches: 10
15 - !operator_fusion {}
16
```

Live Demo